



TUBA TOPRAK

161044116 -Hw5

SYSTEM PROGRAMMING

It uses a producer-consumer model to copy files from a source directory to a destination directory, with multiple consumer threads performing the copying operation concurrently.

The necessary header files and libraries are included.

Global variables are declared, including `buffer_size` (to control the number of threads in the thread pool), `num_consumer_threads` (number of consumer threads), `thread_pool` (an array to store thread IDs), `thread_availability` (an array to track the availability of threads), and `main_thread` (ID of the main thread).

Mutex variables `buffer_mutex` and `output_mutex` are defined to synchronize access to shared resources.

The `Files` struct is defined to hold information about the source and destination directories, as well as the availability of a thread.

The `err_sys` function is a helper function to handle system errors.

The `get_elapsed_time` function calculates the elapsed time between two struct `timeval` timestamps.

The `ignore_sigint` function is a signal handler for `SIGINT` (interrupt signal, usually triggered by pressing Ctrl+C). It cancels all consumer threads, frees memory, and exits gracefully.

The `delete_directory` function is a recursive function to delete a directory and its contents.

The `customer` function is responsible for copying a file from the source directory to the destination directory.

The `customer_thread` function is the entry point for consumer threads. It calls the `customer` function to copy a file and updates the thread availability status.

The `producer` function is the main logic for traversing the source directory and spawning consumer threads to copy files.

The `producer_thread` function is the entry point for the producer thread. It invokes the `producer` function and waits for all consumer threads to complete.

In the `main` function, command-line arguments are parsed to set the buffer size, number of consumer threads, source directory, and destination directory.

The `ignore_sigint` function is registered as the signal handler for `SIGINT`.

Validation is performed on the command-line arguments to ensure positive integer values for the buffer size and number of consumer threads.

Memory is allocated for the thread availability and thread pool arrays.

The thread availability array is initialized to 0 (indicating availability).

The destination directory is deleted (if it exists) and created anew.

Timing is started using `gettimeofday()`.

The buffer mutex and output mutex are initialized.

The main thread creates the producer thread, which starts the file copying process.

The main thread waits for the producer thread to complete.

The mutexes are destroyed. Timing is ended, and the elapsed time is calculated and displayed. Memory is freed.

multithreading is used to improve the performance of the file copying process by leveraging concurrent execution. The program utilizes a producer-consumer model, where the producer thread traverses the source directory and spawns multiple consumer threads to perform the file copying operation.

Here's how the multithreading is implemented:

**Thread Pool:** The program creates a pool of consumer threads, specified by the `num_consumer_threads` variable. The thread pool is represented by the `thread_pool` array, which stores the thread IDs.

**Thread Availability:** To manage the availability of consumer threads, the `thread_availability` array is used. Each element of the array corresponds to a consumer thread, and a value of 0 indicates that the thread is available for processing a file. Initially, all threads are marked as unavailable (0).

**Buffer Mutex:** The `buffer_mutex` is a mutex (mutual exclusion) variable used to synchronize access to the `thread_availability` array. It ensures that only one thread at a time can check and update the availability status.

**Output Mutex:** The `output_mutex` is another mutex variable used to synchronize access to the output, such as printing the copied file information. It ensures that multiple threads don't interfere with each other when writing to the output.

**Producer Thread:** The main thread acts as the producer thread. It invokes the producer function, which recursively traverses the source directory. When it encounters a file, it checks for an available consumer thread using the `thread_availability` array. If an available thread is found, it marks it as unavailable, creates a structure (Files) to hold the file information, and passes it to the consumer thread using the `pthread_create` function. The producer thread then waits for all consumer threads to complete using `pthread_join`.

**Consumer Threads:** The consumer threads are responsible for copying files from the source directory to the destination directory. Each consumer thread executes the `customer_thread` function. Inside this function, the customer function is called to perform the file copy operation. Access to the shared resources (thread availability and output) is protected using mutex locks (`pthread_mutex_lock` and `pthread_mutex_unlock`).

## Controls:

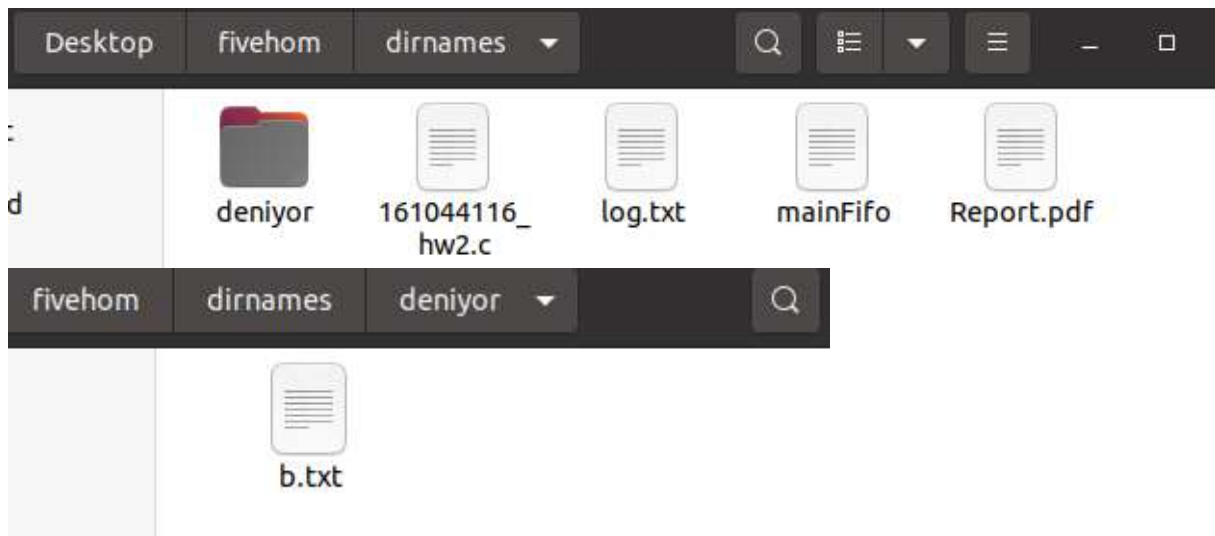
```
toprak@toprak-virtual-machine:~/Desktop/flvehon$ valgrind --leak-check=full --track-origins=yes ./pCp
==23160== Memcheck, a memory error detector
==23160== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==23160== Using Valgrind-3.15.0 and LibVEX; rerun with -h for copyright info
==23160== Command: ./pCp
==23160==
Usage: ./pCp <buffer_size> <num_consumer_threads> <source_directory> <destination_directory>
: Success
==23160==
==23160== HEAP SUMMARY:
==23160==    in use at exit: 0 bytes in 0 blocks
==23160==   total heap usage: 2 allocs, 2 frees, 1,496 bytes allocated
==23160==
==23160== All heap blocks were freed -- no leaks are possible
==23160==
==23160== For lists of detected and suppressed errors, rerun with: -s
==23160== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
toprak@toprak-virtual-machine:~/Desktop/flvehon$
```

```
toprak@toprak-virtual-machine:~/Desktop/fivehom$ ps -aux | grep "Z"
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
toprak    23165  0.0  0.0  11776   652 pts/0    S+   22:47   0:00 grep --color=auto Z
```

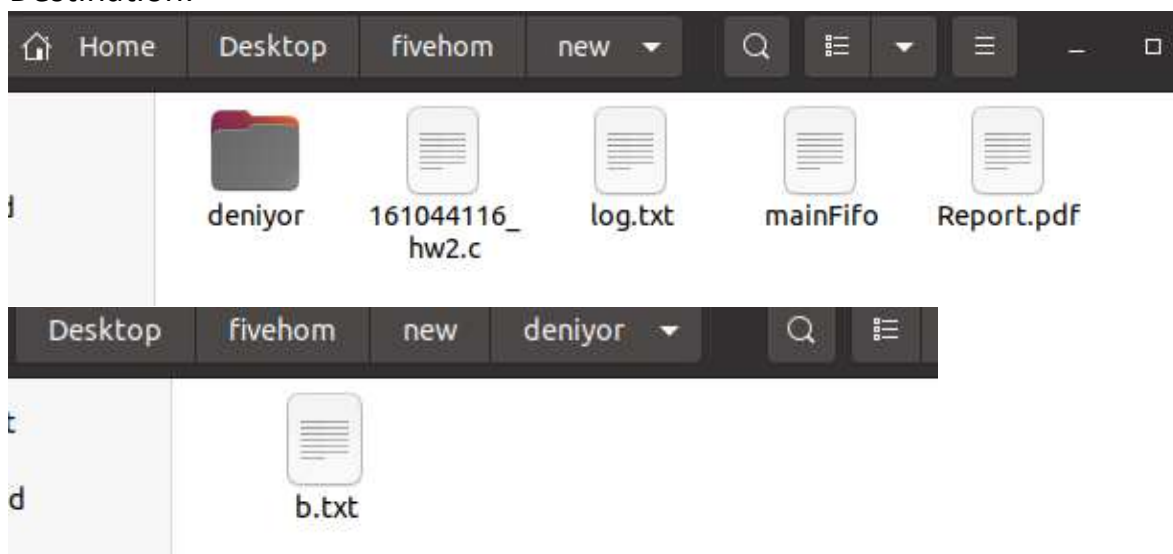
## Output:

```
toprak@toprak-virtual-machine:~/Desktop/fivehom$ ./pCp 100 10 dirnames new
Copying from dirnames folder to new folder.
Copied Folder: new/deniyor
Copied File: new/deniyor/b.txt
Copied File: new/log.txt
Copied Fifo: new/mainFifo
Copied File: new/Report.pdf
Copied File: new/161044116_hw2.c
Elapsed time: 0.002077 seconds
Finished..
toprak@toprak-virtual-machine:~/Desktop/fivehom$
```

## Source:



## Destination:



**Observations:**

As the buffer size increased, the elapsed time decreased. This indicates that a larger buffer allows for more parallelism and reduces the time spent waiting for available consumer threads.

Similarly, increasing the number of consumer threads resulted in decreased elapsed time. More threads allow for higher concurrency, leading to faster file copying.

**Conclusion:**

The multithreaded file copying utility demonstrated improved performance compared to a single-threaded approach. The experiments revealed that larger buffer sizes and a higher number of consumer threads yielded better results in terms of elapsed time.